Ryan Carlson (rcarlson)
Software Engineering – Homework 3
October 7, 2013

I'll take each task in turn, discussing how my approach.

Task 1: CPE

The CPE was fairly straightforward. I copied over the `FileSystemCollectionReader` and `XmiWriterCasConsumer` verbatim from the UIMA examples. For the cas consumer, I modified it to keep a running total of the precision and the number of elements that have been evaluated, and then I implemented the `collectionProcessComplete` to report the average `P@N` score to standard out. Then I just created a CPE that involved all those components with `hw2-rcarlson-aae` as the analysis engine.

Task 2: UIMA-AS

Here I created a custom resource descriptor to contact the server and get the Stanford Core NLP annotations. Then I needed a way to integrate those results into my scoring system. So I created a new descriptor called `StanfordNameEntityScoreAnnotator`. This descriptor takes question / answer pairs and assigns a score based on the overlap of named entities contained in each.

At this point, we also need a way to combine scores from multiple sources (Stanford and CosineSimilarity, in my case), so I created a `CollapseAnswerScoreAnnotator` that looks at the `AnswerScore` instances that cover the same span and sums their scores together into a new `AnswerScore` instance. This then gets passed to the evaluator. The evaluator only looks at these collapsed `AnswerScore` instances

The accuracy remained identical to the results from the second assignment. In both cases, the `hw2-rcarlson-aae` step took under a half second (between 300 and 400 milliseconds). The `scnlp-rcarlson-client` step took just over one second (about 1250 milliseconds). The entire CPE takes 2 seconds, which is about one second slower than without the Stanford annotations.

Note that I modified `hw3-rcarlson-CPE` to add the call to the Stanford service.

The last step was fairly straightforward once `CLASSPATH` issues were sorted out. The deployment descriptor for `hw2-rcarlson-aae` was created, and then I made a client (`hw2-rcarlson-client`) to point to the service with the queue name of

"rcarlson-hw2-aae-client-queue." Then I just created a CPE that uses the service, started the broker, deployed the service, and ran the CPE.

Note that I couldn't figure out how to use relative paths, so I only have absolute paths in my CPEs. This might cause a problem for automatic testing and whatnot. Also, if it's interesting, my CLASSPATH was set to

```
/Users/rcarlson/local/eclipse-workspaces/software-
engineering/cmu11791-hw3-rcarlson/hw3-
rcarlson/target/dependency/:/Users/rcarlson/local/eclips
e-workspaces/software-engineering/cmu11791-hw3-
rcarlson/hw3-rcarlson/target/classes/
```

I highlighted the sort of "relative" parts in bold.